

## Background and Motivation

Major issues in mainstream adoption of FPGAs:

- Difficulty of accelerator design at low level
- Long compilation times (Place and route)
- Poor design productivity

One possible solution is to use FPGA Overlay:

- Accelerator design in a high level language
- Fast compilation and development cycles
- Easy to use even by novice programmers
- Puts the FPGA in the hands of software developers.

Another solution is to use novel high level synthesis tools:

- HLS tools to generate C to RTL, eg: Vivado HLS, SDSoC
- OpenCL to hardware design synthesis eg: AOCL, SDAccel

## Contributions

- Analysis of Overlay architectures like Vectorblox MXP and Altera OpenCL SDK as alternatives to pure RTL design flow. Highlighting the ease of use and fast learning curves of these methods.
- Benchmarking and comparison of timing performance as well as operations per cycle for the Vectorblox MXP processor and the Altera OpenCL Implementation with arm CPU implementation.

## Observations

The Vectorblox MXP soft vector processor:

- Extremely short learning curve
- The Vectorblox MXP C/C++ api is extremely easy to use.
- Full control of DMA and execution to programmer.
- Compilation and debugging time equal to that of traditional C/C++ debugging.

The Altera OpenCL SDK:

- Minimal modification of traditional OpenCL code required.
- kernel optimizations like vectorization, hardware replication for parallelism.
- Dynamic re-programming of FPGA at runtime with totally new kernel.
- Host code portability across all devices.



## Conclusions and Future Work

- More popularity of overlays and high level synthesis tools.
- Place and route efficiency improvements in OpenCL to hardware.
- Overall generated hardware efficiency improvement.
- More awareness in the software world and amongst software developers.

## Efficient Overlay Architecture

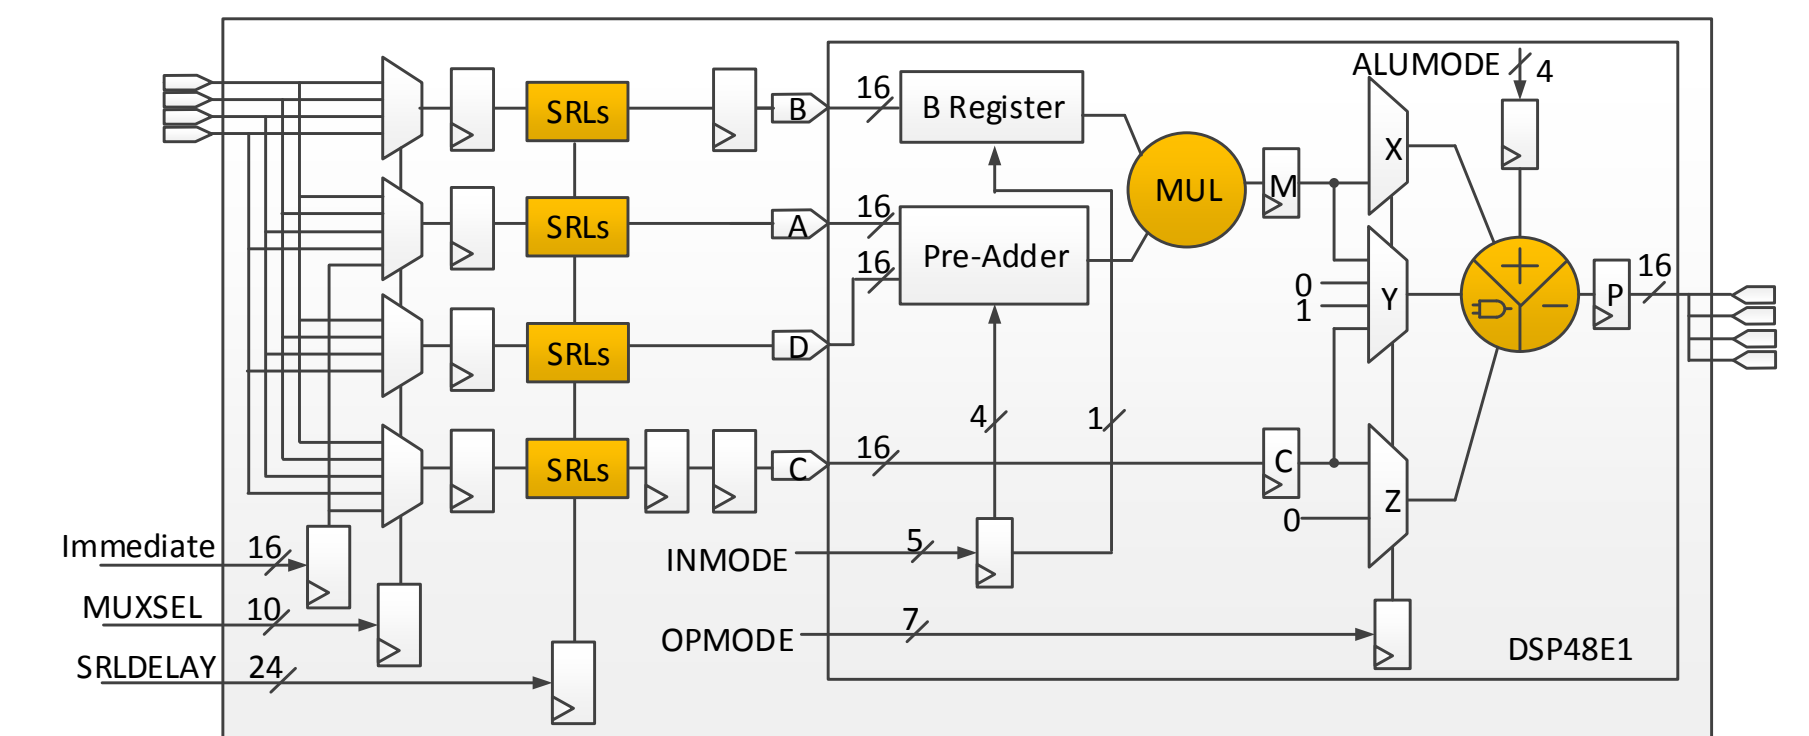
2D array of tiles:

- Programmable functional unit (FU) and routing resources in each tile
- Functional units interconnected via an island-style routing network
- Coarse grained switch boxes, connection boxes and routing channels as programmable routing resources
- Customizable channel width (CW), number of tracks in a routing channel



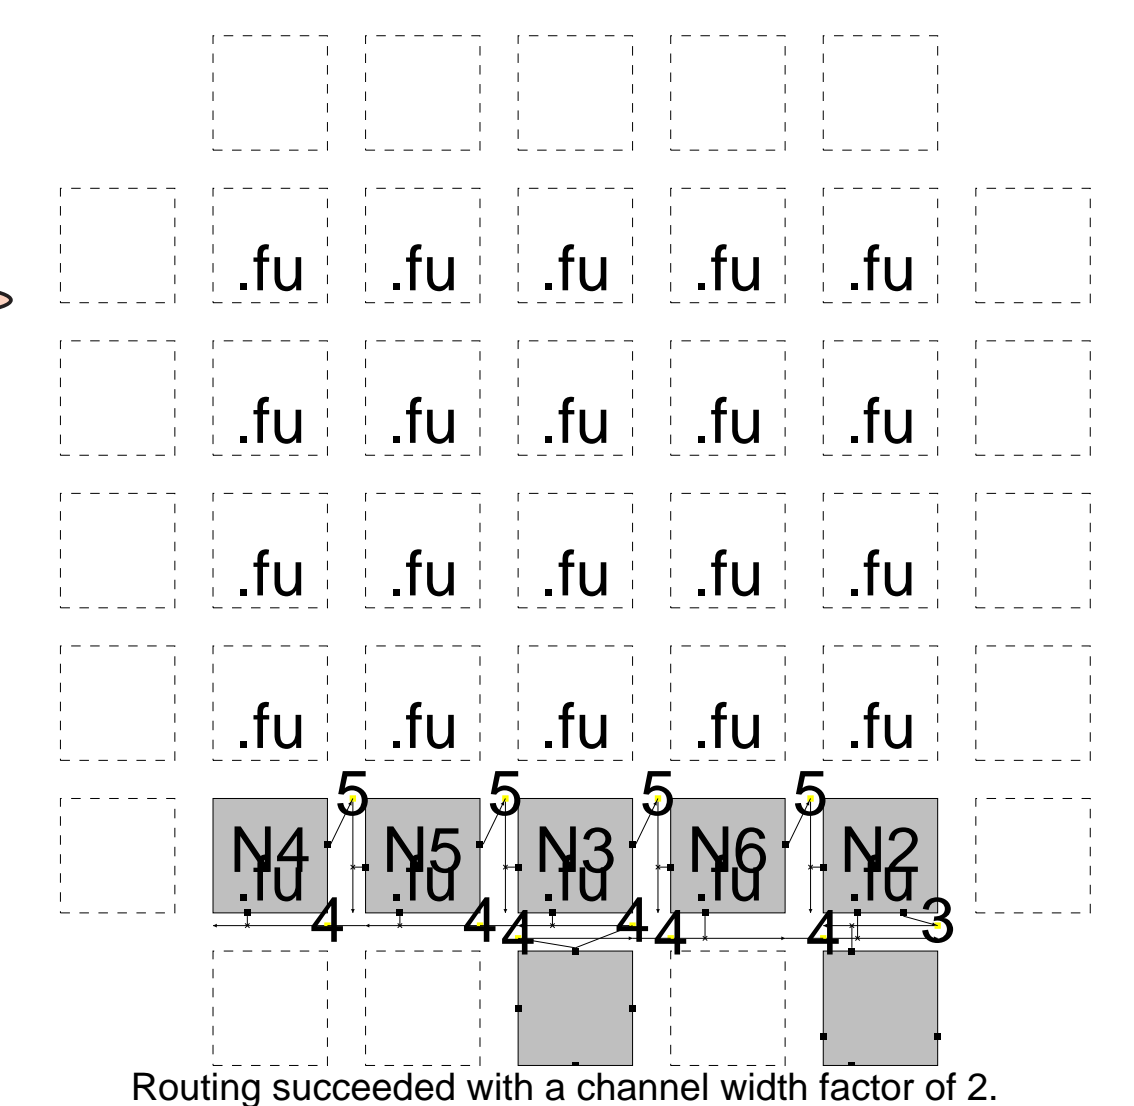
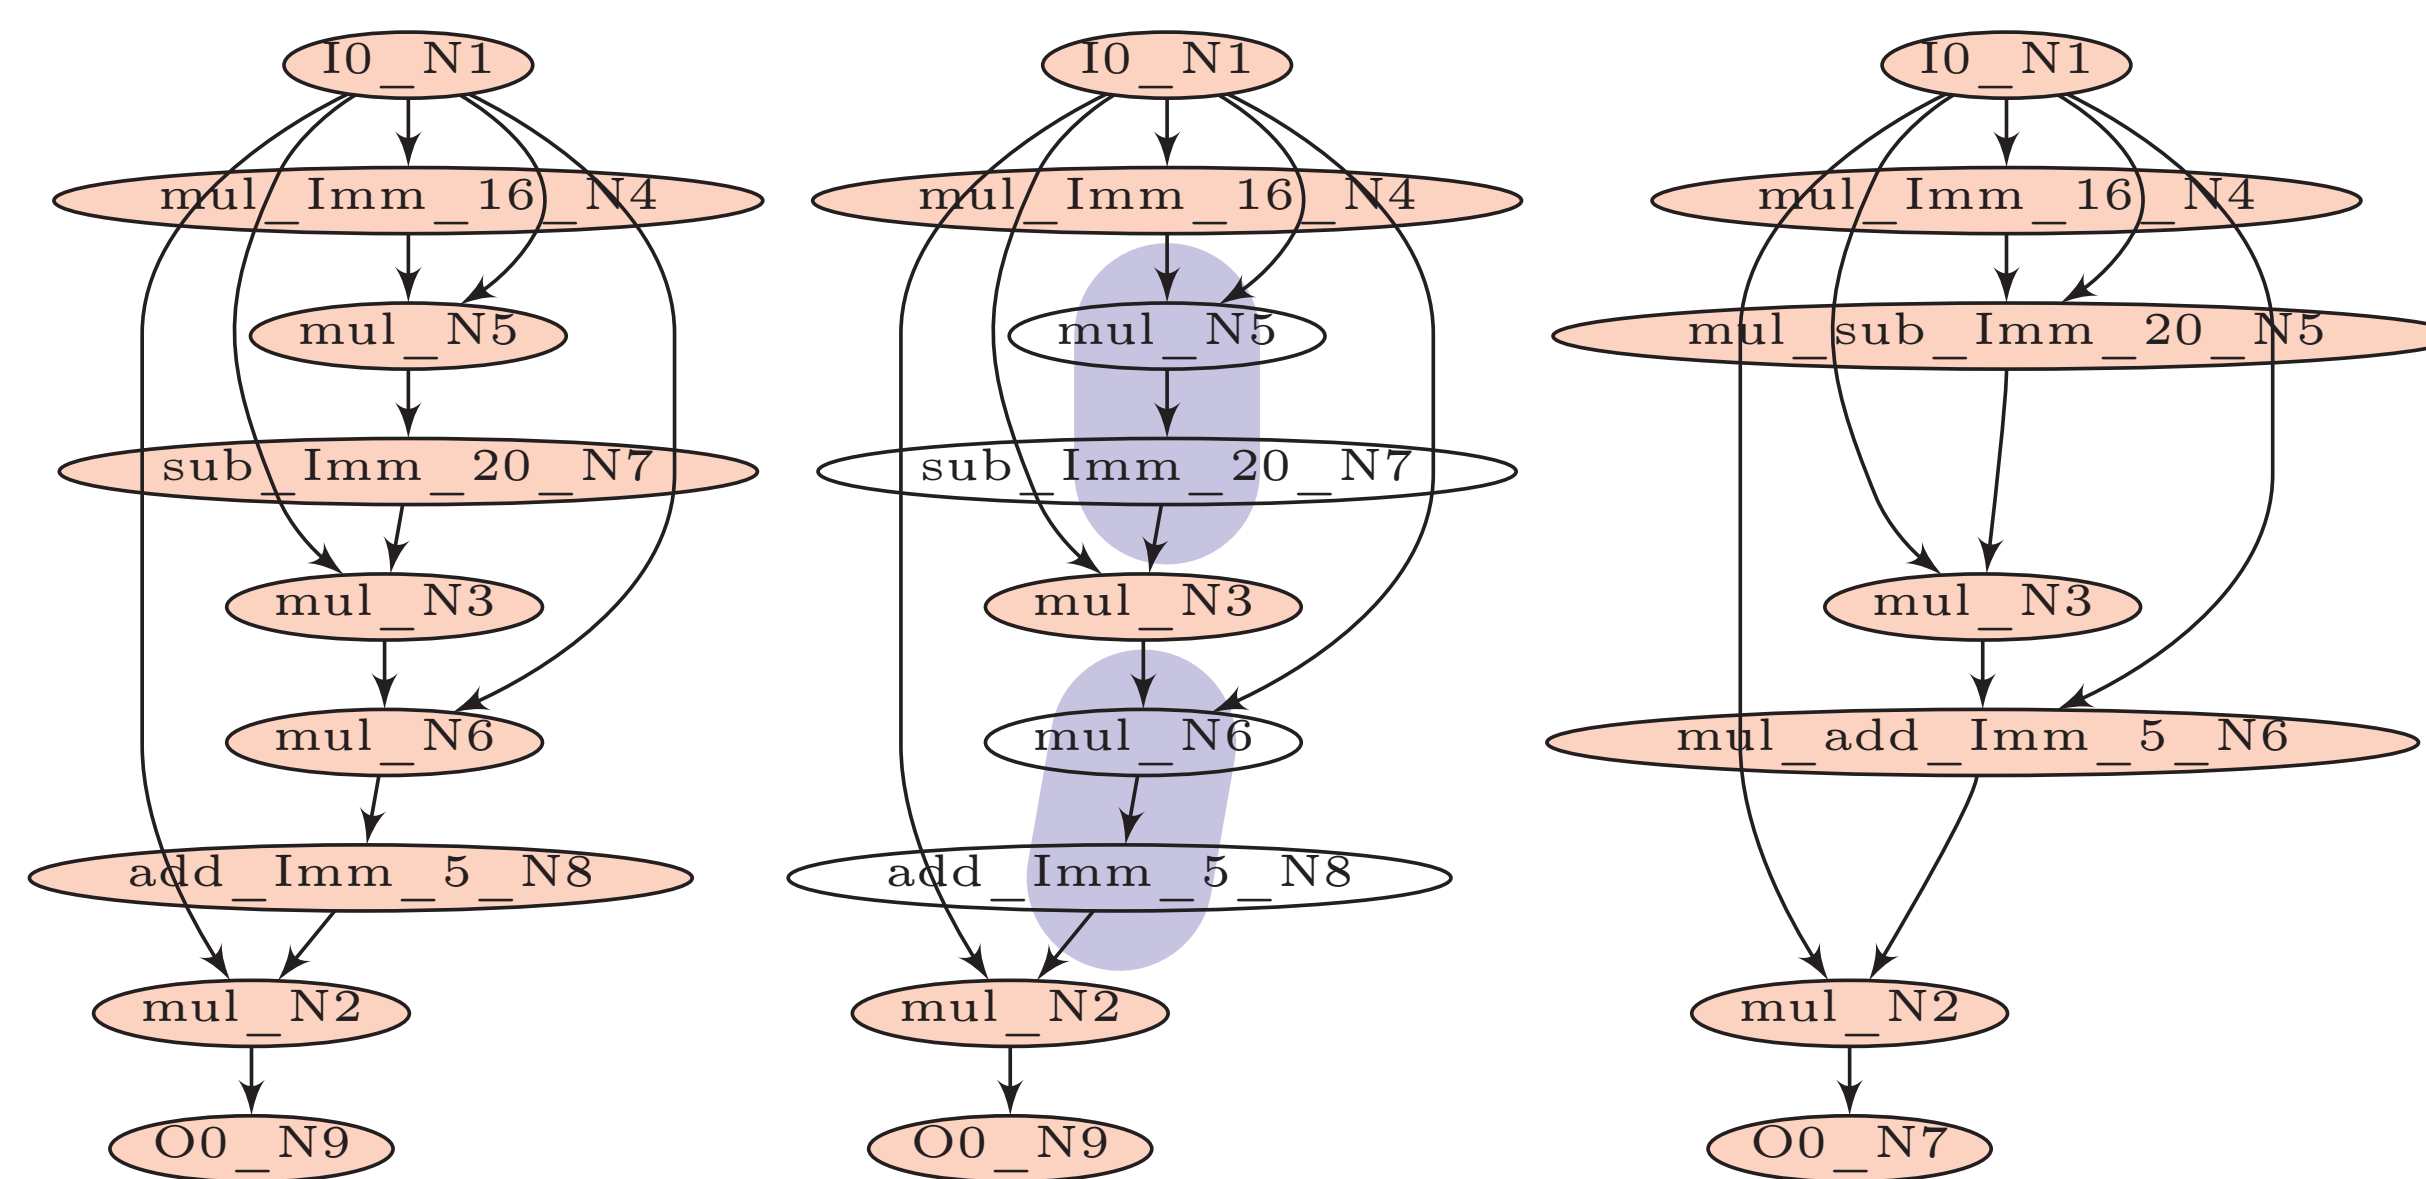
## DSP Block Based Functional Unit

- Fully pipelined DSP48E1 as a programmable processing element (PE)
- Achievable frequency near theoretical limits for providing high throughput
- A pre-adder, a multiplier and an ALU inside the functional unit
- Can support upto 3 operations
- MUX based reordering logic to handle logical inequivalence at the PE inputs
- SRL based variable-length shift registers for balancing pipeline latencies






## Rapid, Vendor-Independent, Automated Mapping of Compute Kernels

- C to DFG Transformation: DFG generation from a C description of the compute kernel
- DSP48E1 Aware Mapping: Compute node merging based on the capability of the DSP block
- Placement and Routing of FU Netlist: Using VPR for mapping nodes in the graph to the DSP blocks, and edges onto the coarse grained tracks
- Latency Balancing: Parsing VPR output files and generating a routing resource graph to determine the latency imbalance at each node and hence the required delays at the FU inputs



## Experimental Evaluation

- Two example overlays on Zynq device to execute the benchmark set: a 5x5 Overlay-I with CW=2 operating at 370 MHz and a 7x7 Overlay-II with CW=4 operating at 300 MHz
- RTL generation of benchmarks using Vivado HLS for performance (throughput) comparison



Benchmark	Benchmark Characteristics				Routability		Overlay Results			HLS Implementation Results				
	i/o nodes	op nodes	merged nodes	savings	CW=2	CW=4	Latency	MLI	GOPS	Latency	Fmax	GOPS	Slices	DSPs
chebyshev	1/1	7	5	28%	3x3	3x3	49	36	<b>2.59</b>	13	333	<b>2.3</b>	24	3
sgfilter	2/1	18	10	44%	4x4	4x4	54	31	<b>6.66</b>	11	278	<b>5.0</b>	40	8
mibench	3/1	13	6	53%	3x3	3x3	47	35	<b>4.81</b>	9	295	<b>3.8</b>	81	3
qspline	7/1	26	22	15%	5x5	5x5	76	64	<b>9.62</b>	21	244	<b>6.3</b>	126	14
poly1	2/1	9	6	33%	3x3	3x3	34	22	<b>3.33</b>	12	285	<b>2.66</b>	62	4
poly2	2/1	9	6	33%	3x3	3x3	29	7	<b>3.33</b>	11	295	<b>2.65</b>	45	4
poly3	6/1	11	7	36%	3x3	3x3	31	11	<b>4.07</b>	12	250	<b>2.75</b>	52	6
poly4	5/1	6	3	50%	2x2	2x2	24	12	<b>2.22</b>	7	312	<b>1.87</b>	36	3
atax	12/3	60	36	40%	—	6x6	72	58	<b>18.0</b>	13	263	<b>15.8</b>	78	18
bicg	15/6	30	18	40%	—	6x6	46	32	<b>9.0</b>	7	270	<b>8.1</b>	91	18
trmm	18/9	54	36	33%	—	7x7	58	30	<b>16.2</b>	8	222	<b>11.9</b>	105	36
syrk	18/9	72	45	37%	—	7x7	41	19	<b>21.6</b>	10	250	<b>18</b>	237	24